

✓ A Data Analysis of Aircraft Incidents for Business Strategy Minimizing Risk

Overview

In this project, I analyzed aviation accident data from the National Transportation Safety Board (NTSB) to help guide safer aircraft purchases. Using Python and pandas, I cleaned and explored the data, focusing on trends in manufacturer safety, injury severity, and number of engines. The analysis showed that aircraft from certain manufacturers and those with more engines tend to be involved in fewer and less severe accidents. These insights informed three business recommendations to support data-driven decisions as the company enters the aviation industry.

Business Understanding

The business problem focuses on minimizing risk as the company expands into aviation by purchasing aircraft for commercial and private use. To support safe investment decisions, I analyzed historical accident data to identify manufacturers with fewer incidents, aircraft with more engines, and models linked to less severe outcomes. These factors directly impact safety, operational costs, and reputation, helping the company make informed, low-risk purchasing choices.

✓ Data Understanding

The dataset comes from the NTSB and contains records of civil aviation accidents from 1962 to 2023, including aircraft manufacturer, model, number of engines, location, and injury severity. It reflects the safety history of different aircraft, which is critical for evaluating investment risks in aviation. The analysis focuses on accident severity as the target variable, using both categorical and numerical data to uncover safety trends.

```
# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset
df = pd.read_csv('/content/AviationData.csv', encoding='latin1')

# Preview first few rows
df.head()
```



	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude	Airport.Code	Airport.Name	...	Purpose.of.flight	Air.carrier	Total.Fatal
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States	NaN	NaN	NaN	NaN	...	Personal	NaN	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States	NaN	NaN	NaN	NaN	...	Personal	NaN	
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States	36.922223	-81.878056	NaN	NaN	...	Personal	NaN	
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States	NaN	NaN	NaN	NaN	...	Personal	NaN	
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States	NaN	NaN	NaN	NaN	...	Personal	NaN	

5 rows × 31 columns



The dataset contains accident reports from 1962 to 2023, including fields like aircraft manufacturer, model, number of engines in the aircraft, and injury severity. The target focus is on accident severity and safety trends among manufacturers and aircraft types.

Basic information about the dataset

```
print("\nData Information:")
df.info()
```



```
Data Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29623 entries, 0 to 29622
Data columns (total 31 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Event.Id              29623 non-null object
1   Investigation.Type     29623 non-null object
2   Accident.Number       29623 non-null object
3   Event.Date            29623 non-null object
4   Location              29613 non-null object
5   Country               29482 non-null object
6   Latitude              8 non-null     float64
7   Longitude             8 non-null     float64
8   Airport.Code          15658 non-null object
9   Airport.Name          17379 non-null object
10  Injury.Severity        29623 non-null object
11  Aircraft.damage        28978 non-null object
12  Aircraft.Category      3665 non-null  object
13  Registration.Number    29612 non-null object
14  Make                  29616 non-null object
15  Model                 29606 non-null object
16  Amateur.Built         29622 non-null object
17  Number.of.Engines      29291 non-null float64
18  Engine.Type           29619 non-null object
19  FAR.Description        3665 non-null  object
```

```
20 Schedule                4849 non-null object
21 Purpose.of.flight       29579 non-null object
22 Air.carrier              1584 non-null object
23 Total.Fatal.Injuries     29478 non-null float64
24 Total.Serious.Injuries   29432 non-null float64
25 Total.Minor.Injuries     29428 non-null float64
26 Total.Uninjured          29496 non-null float64
27 Weather.Condition        29621 non-null object
28 Broad.phase.of.flight    29387 non-null object
29 Report.Status            29623 non-null object
30 Publication.Date          17176 non-null object
dtypes: float64(7), object(24)
memory usage: 7.0+ MB
```

Displaying the number of rows and columns

```
print("Shape of the dataset:", df.shape)
```

➡ Shape of the dataset: (29623, 31)

Statistical Summary

```
print("\nSummary Statistics:")
df.describe(include='all')
```

➡

Summary Statistics:														
	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country	Latitude	Longitude	Airport.Code	Airport.Name	...	Purpose.of.flight	Air.carrier	Total
count	29623	29623	29623	29623	29613	29482	8.000000	8.000000	15658	17379	...	29579	1584	
unique	29171	2	29623	3636	10627	59	NaN	NaN	5003	8905	...	12	1251	
top	20001214X45071	Accident	ANC92LA021	1982-05-16	ANCHORAGE, AK	United States	NaN	NaN	NONE	PRIVATE	...	Personal	United Airlines	
freq	3	28583	1	25	211	29172	NaN	NaN	617	102	...	16713	33	
mean	NaN	NaN	NaN	NaN	NaN	NaN	44.914271	-110.695764	NaN	NaN	...	NaN	NaN	
std	NaN	NaN	NaN	NaN	NaN	NaN	11.772969	36.243789	NaN	NaN	...	NaN	NaN	
min	NaN	NaN	NaN	NaN	NaN	NaN	30.757778	-173.240000	NaN	NaN	...	NaN	NaN	
25%	NaN	NaN	NaN	NaN	NaN	NaN	38.135556	-128.370625	NaN	NaN	...	NaN	NaN	
50%	NaN	NaN	NaN	NaN	NaN	NaN	44.243194	-101.121527	NaN	NaN	...	NaN	NaN	
75%	NaN	NaN	NaN	NaN	NaN	NaN	46.645833	-84.717222	NaN	NaN	...	NaN	NaN	
max	NaN	NaN	NaN	NaN	NaN	NaN	70.333333	-70.758333	NaN	NaN	...	NaN	NaN	


11 rows × 31 columns

Data Preparation

I cleaned the dataset by dropping irrelevant columns, handling missing values through removal or forward-filling, and standardizing categorical data. These steps ensured the dataset was reliable and focused on key factors related to aircraft safety and risk assessment.

Checking for missing values

```
missing_values = df.isnull().sum()
missing_values = missing_values[missing_values > 0]
print("\nMissing Values:\n", missing_values)
```



Missing Values:	
Location	10
Country	141
Latitude	29615
Longitude	29615
Airport.Code	13965
Airport.Name	12244
Aircraft.damage	645
Aircraft.Category	25958
Registration.Number	11
Make	7
Model	17
Amateur.Built	1
Number.of.Engines	332
Engine.Type	4
FAR.Description	25958
Schedule	24774
Purpose.of.flight	44
Air.carrier	28039
Total.Fatal.Injuries	145
Total.Serious.Injuries	191
Total.Minor.Injuries	195
Total.Uninjured	127
Weather.Condition	2
Broad.phase.of.flight	236
Publication.Date	12447
dtype: int64	

Dropping columns with more than 30% missing values

```
threshold = len(df) * 0.3
cols_to_drop = missing_values[missing_values > threshold].index
df = df.drop(columns=cols_to_drop, errors='ignore')
```

Dropping irrelevant columns manually

```
df = df.drop(['Location', 'Country', 'Publication_Date', 'Report_Status',
             'Aircraft_Category', 'Broad_phase_of_flight', 'Schedule', 'Air_carrier',
             'FAR_Description', 'Longitude', 'Latitude', 'Airport_Code', 'Airport_Name'], axis=1, errors='ignore')
df.head()
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Injury.Severity	Aircraft.damage	Registration.Number	Make	Model	Amateur.Built	Number.ofEngines	Engine.Type	Pur
0	20001218X45444	Accident	SEA87LA080	1948-10-24	Fatal(2)	Destroyed	NC6404	Stinson	108-3	No	1.0	Reciprocating	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	Fatal(4)	Destroyed	N5069P	Piper	PA24-180	No	1.0	Reciprocating	
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Fatal(3)	Destroyed	N5142R	Cessna	172M	No	1.0	Reciprocating	
3	20001218X45448	Accident	LAX96LA321	1977-06-19	Fatal(2)	Destroyed	N1168J	Rockwell	112	No	1.0	Reciprocating	
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Fatal(1)	Destroyed	N15NY	Cessna	501	No	NaN	NaN	

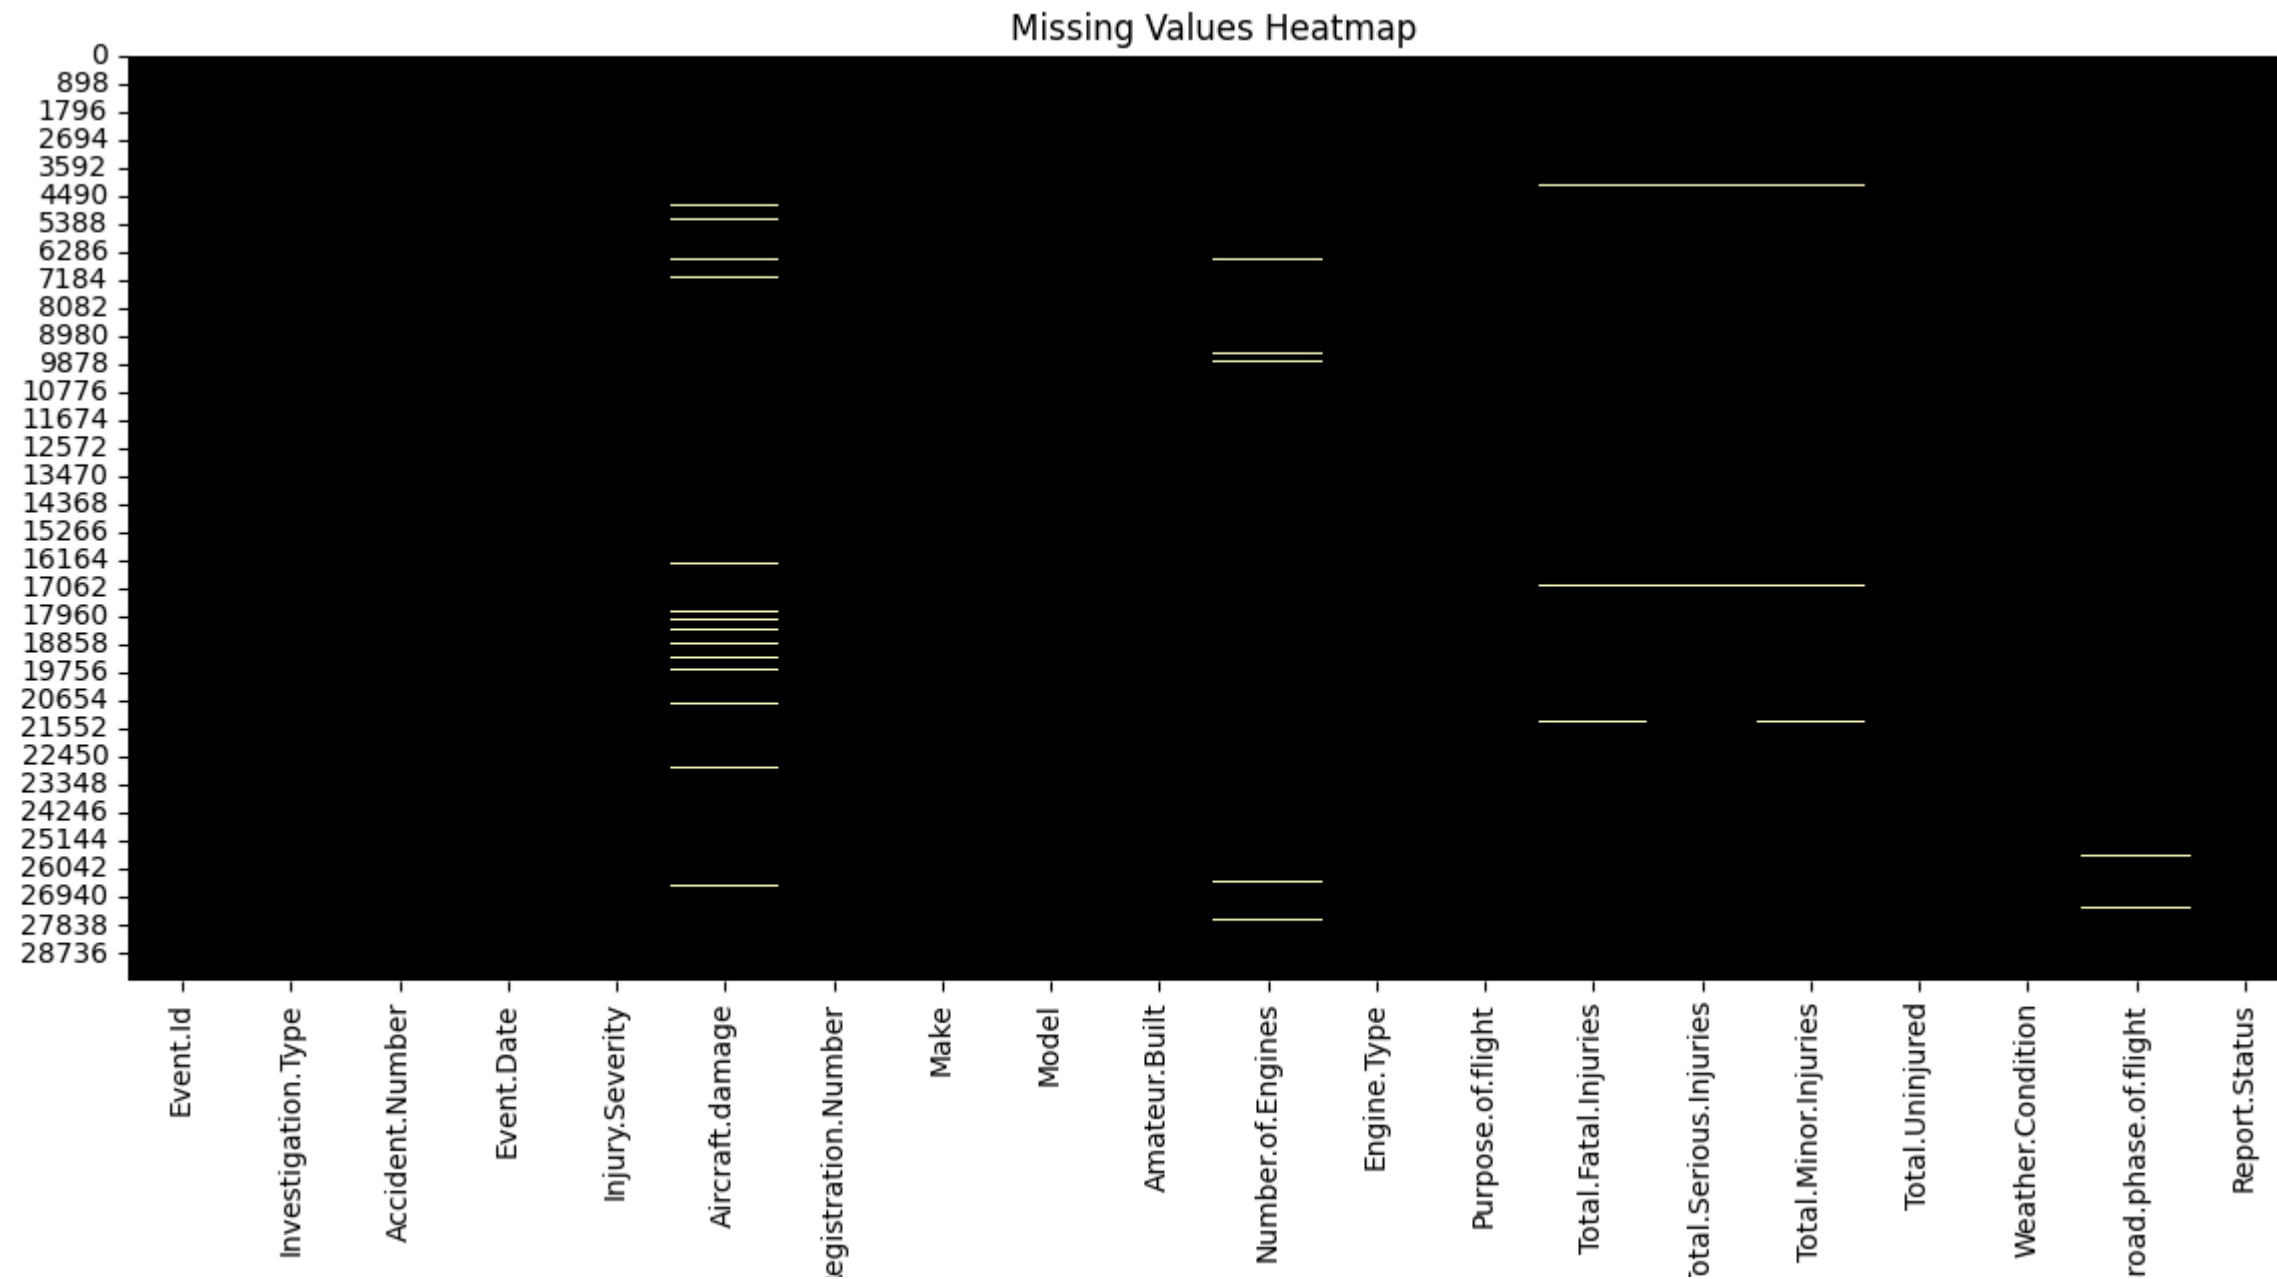
Next steps:

View recommended plots

New interactive sheet

Visualizing missing data

```
plt.figure(figsize=(14,6))
sns.heatmap(df.isnull(), cbar=False, cmap='magma')
plt.title('Missing Values Heatmap')
plt.show()
```



▼ Filling remaining values with forward fill

```
data = df.fillna(method='ffill')
```



```
<ipython-input-9-8c21be8192a7>:1: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
data = df.fillna(method='ffill')
```

▼ Verifying missing after cleaning

```
print("\nMissing Values after cleaning:\n", data.isnull().sum().sum())
```



```
Missing Values after cleaning:
0
```

✓ Check final data shape

```
df.shape
```

```
↔ (29623, 20)
```

I dropped irrelevant columns and handled missing data appropriately to maintain focus on aircraft safety factors, ensuring a clean dataset for analysis.

✓ Exploratory Data Analysis(EDA)

I used exploratory data analysis (EDA) with groupings, aggregations, and visualizations to uncover trends related to aircraft safety. This descriptive approach was appropriate because the project focused on identifying safer aircraft models, not predictive modeling.


✓ Top 10 Aircraft Manufacturers

✓ Analyze aircraft manufacturer accident counts

```
manufacturer_counts = df['Make'].value_counts().head(10)
```

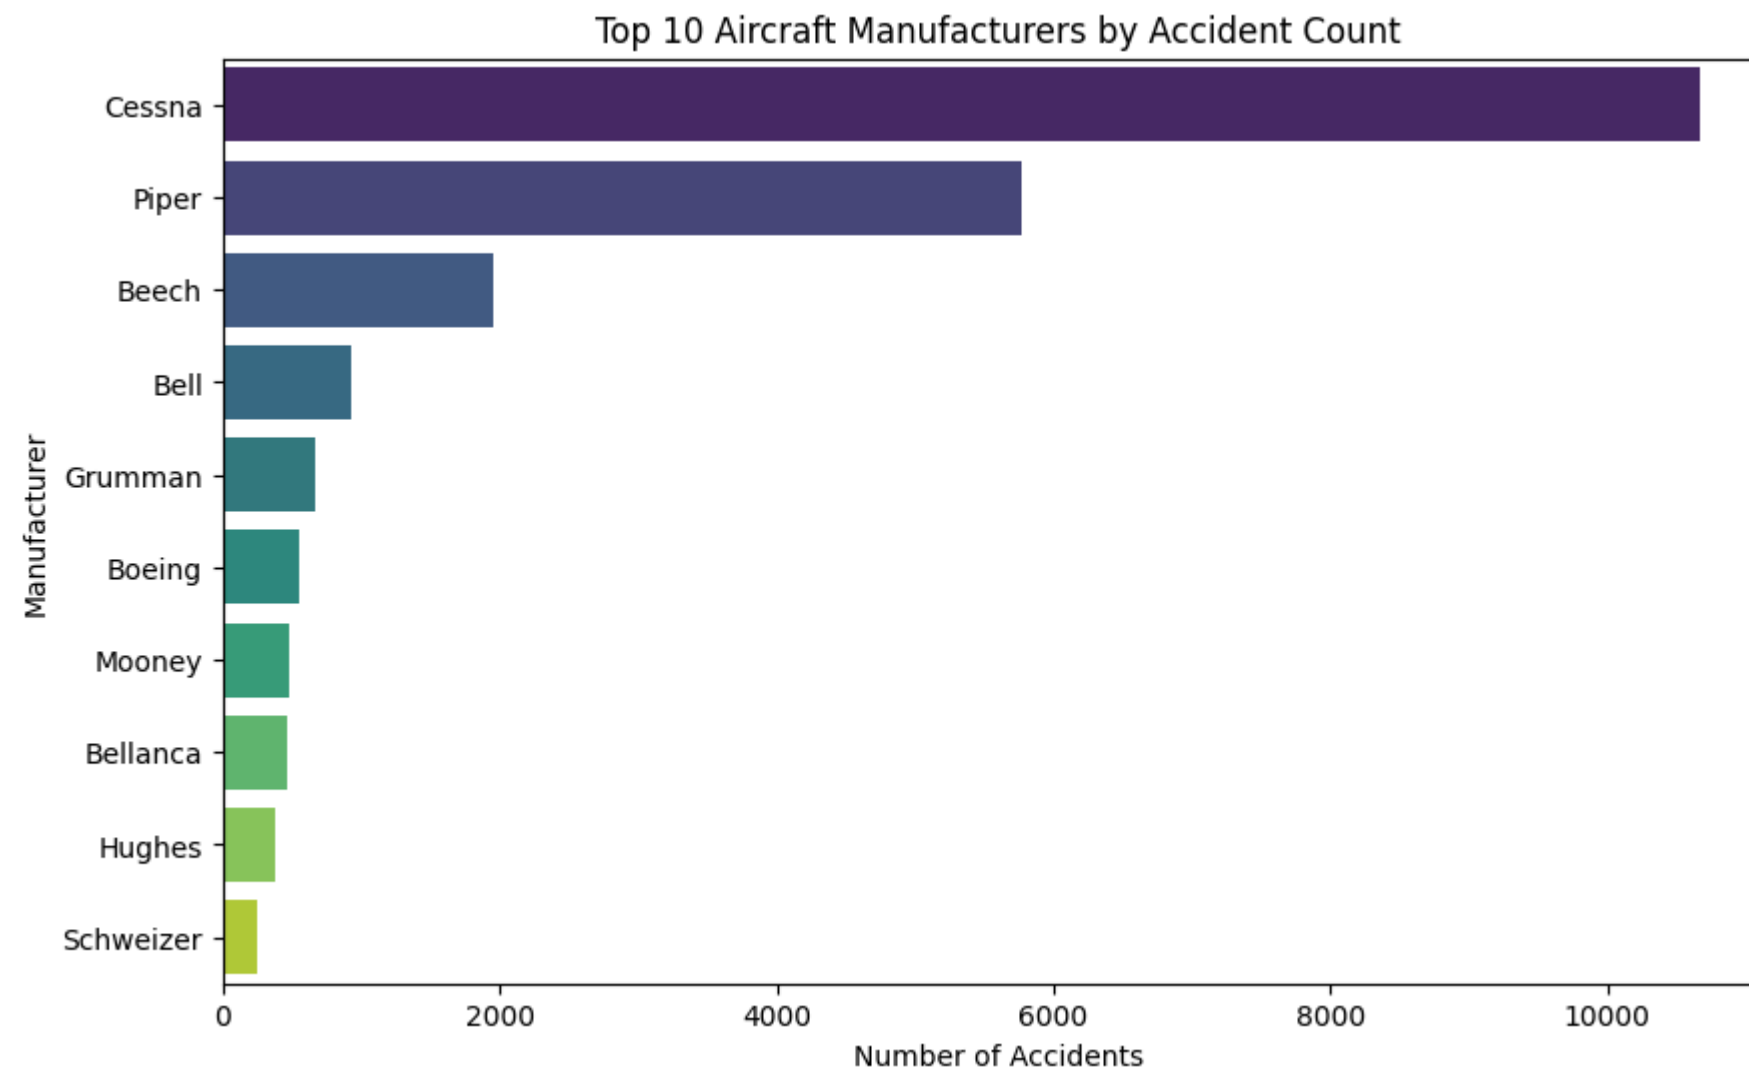


```
plt.figure(figsize=(10,6))
sns.barplot(x=manufacturer_counts.values, y=manufacturer_counts.index, palette='viridis')
plt.title('Top 10 Aircraft Manufacturers by Accident Count')
plt.xlabel('Number of Accidents')
plt.ylabel('Manufacturer')
plt.show()
```

 <ipython-input-13-b9eddfef0f59>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=manufacturer_counts.values, y=manufacturer_counts.index, palette='viridis')
```




Accidents by the number of engines in the aircraft

✓ Analyze the number of engines by the aircraft

```
engine_data = df[['Number.of.Engines', 'Total.Fatal.Injuries']]
engine_accident_counts = engine_data.groupby('Number.of.Engines').size().reset_index(name='Accident_Count')
engine_accident_counts = engine_accident_counts.sort_values(by='Accident_Count', ascending=False)
engine_accident_counts['Engine_Risk_Level'] = engine_accident_counts['Accident_Count'].apply(lambda x: 'High Risk' if x > 10 else 'Low Risk')
```


```
print(engine_accident_counts)
```



	Number.of.Engines	Accident_Count	Engine_Risk_Level
1	1.0	24485	High Risk
2	2.0	3865	High Risk

0	0.0	474	High Risk
3	3.0	283	High Risk
4	4.0	184	High Risk

```
plt.figure(figsize=(10,6))
plt.bar(engine_accident_counts['Number.ofEngines'].astype(str), engine_accident_counts['Accident_Count'], color='teal')
plt.xlabel('Number of Engines')
plt.ylabel('Accident Count')
plt.title('Accident Counts by Number of Engines')
plt.show()
```



I conclude that the higher the number of engines the lower the accident count.

✦ Injury Severity Distribution

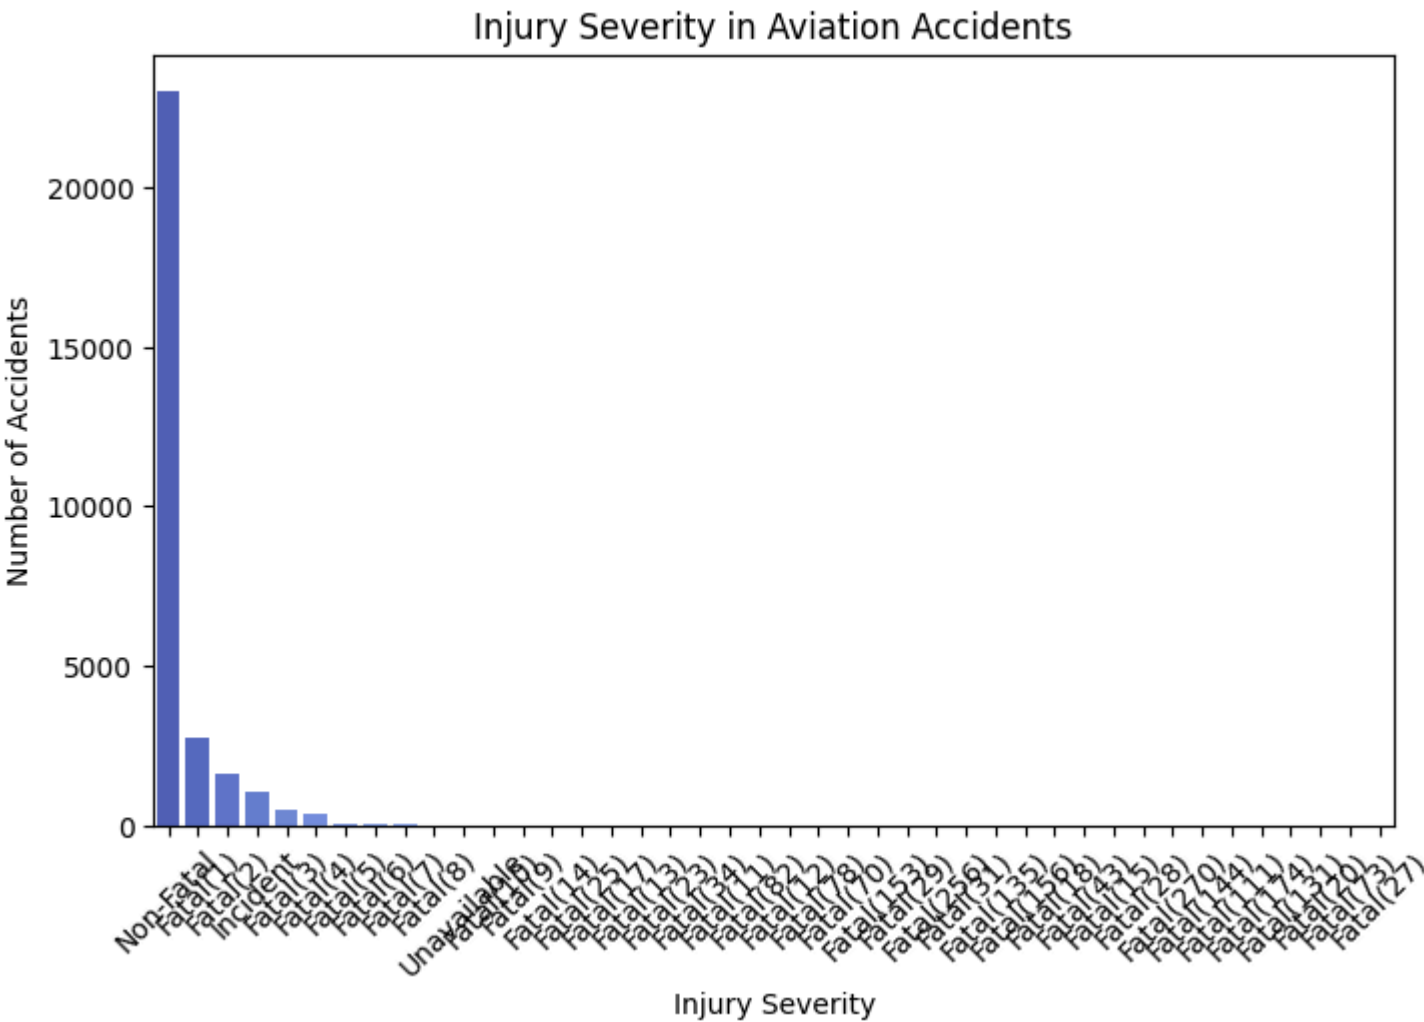
✦ Analyze injury severity

```
injury_counts = df['Injury.Severity'].value_counts()
```

```
plt.figure(figsize=(8,5))
sns.barplot(x=injury_counts.index, y=injury_counts.values, palette='coolwarm')
plt.title('Injury Severity in Aviation Accidents')
plt.xlabel('Injury Severity')
plt.ylabel('Number of Accidents')
plt.xticks(rotation=45)
plt.show()
```

```
<ipython-input-17-68a949998212>:2: FutureWarning:
    Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=injury_counts.index, y=injury_counts.values, palette='coolwarm')
```

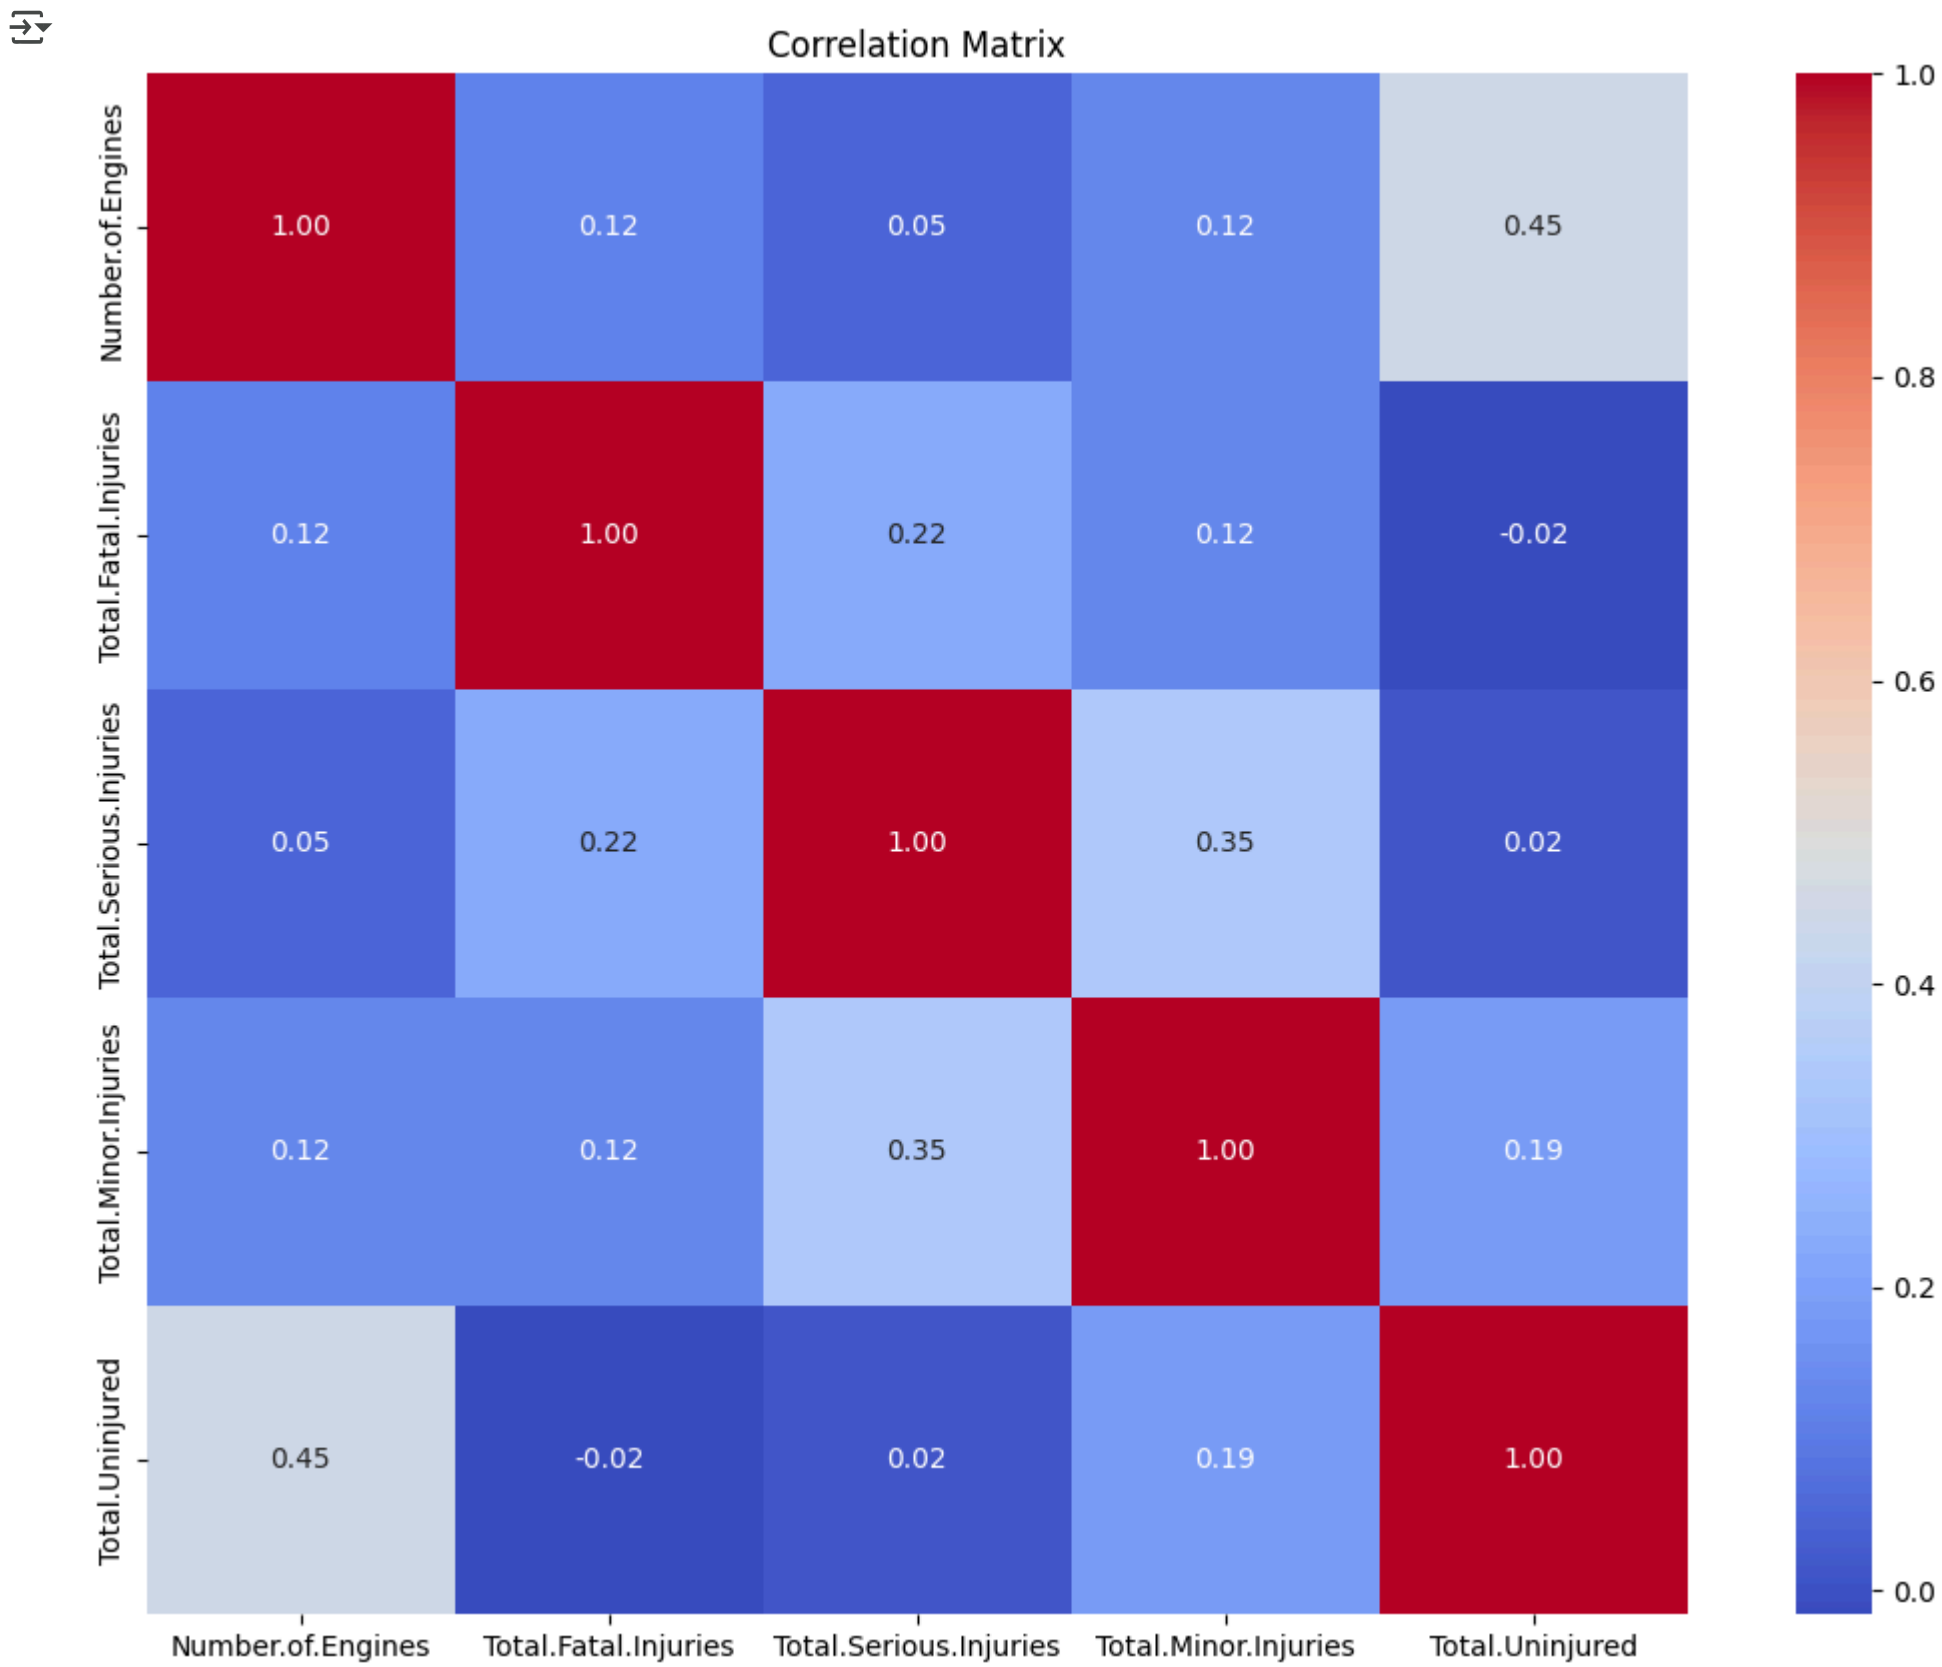


✖ Correlation Matrix

```
# Select only numeric columns
numeric_df = df.select_dtypes(include=[np.number])

if not numeric_df.empty:
    plt.figure(figsize=(14,10))
    sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt='.2f', square=True)
    plt.title('Correlation Matrix')
    plt.show()
```

```
else:  
    print("No numeric data available for correlation heatmap.")
```



Evaluation

The analysis successfully identified safer aircraft manufacturers and revealed that planes with more engines are linked to less severe accidents. Since this was a descriptive analysis, evaluation focused on how well the insights addressed business questions and supported

risk-aware decision-making.

Recommendations

1. Focus investments on aircraft models built by manufacturers who show consistently low numbers of accidents and incidents over time.
- Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.